

Cross-Reader Generality of PAQ-KV: A Multi-VLM Study (Internal Companion Note)

Anonymous submission

Abstract

PAQ-KV is a training-free, index-free, single-reader method: a vision-language model (VLM) prefills a multi-page document once, and per query computes its *own* question-to-context attention, selects the top $\approx 5\%$ of 64-token visual key/value (KV) blocks, and decodes under a reversible 4-D restricted-attention mask over its own KV cache. Its headline result—beating a trained ColQwen2 retriever on MMLongBench-Doc accuracy—is established on Qwen3-VL-8B and confirmed out-of-sample. This note asks the natural follow-up: *does the advantage generalize to other VLM readers?* We port the full method (reader-own scoring, selection, and masked decode) to additional VLMs and report matched-budget results on a fixed 80-document / 444-query development (DEV) pin. Findings so far: on Qwen3-VL-4B (same architecture, half the parameters) the method largely *holds* at byte-identical budget ($\Delta = +0.030$, $LB_{95} = -0.004$, a near-win that denoises above full context); on Qwen2.5-VL-7B (same Qwen-VL family) it is a *statistical dead heat* ($\Delta = +0.007$, $LB_{95} = -0.037$); on InternVL3.5-8B (a different VLM whose language backbone is Qwen3) it *significantly underperforms* ($\Delta = -0.067$, $LB_{95} = -0.106$). Together these characterize a **mechanistic boundary** located at architecture rather than scale: the method’s benefit tracks how good the host reader’s own attention is as a below-page retriever. Three further readers (Molmo-7B-D, MiniCPM-V-4.5, PaliGemma-2) were each found incompatible with the current setup for distinct, honest reasons (remote-code version drift for the first two; an $\approx 8K$ context that cannot hold a multi-page document for the third), so the study’s evidence rests on the four readers above. This is an internal companion note kept separate from the submission; the shared team decides what, if anything, to fold in.

1 Setup

1.1 The method and the established result

PAQ-KV prefills a full multi-page document once into a VLM’s KV cache. For each query it (i) reads the reader’s own question-to-context attention (scorer v0: mean over heads and layers), (ii) partitions the retained visual KV into 64-token stripe blocks and selects the top $\approx 5\%$ by attention

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

score, and (iii) decodes the answer under a reversible 4-D additive attention mask that hides all non-selected visual columns while always keeping the (small) non-visual scaffold. It trains nothing, builds no external index, and uses a single model.

On **Qwen3-VL-8B**, out-of-sample on a sealed 45-document / 244-query MMLongBench-Doc split, PAQ-KV at 5% scores **0.3957** versus a trained ColQwen2 retriever’s **0.3202** ($+0.0755$; doc-cluster $LB_{95} = +0.0290$). On the 80-document / 444-query DEV pin the same beat is $+0.079$ ($LB_{95} = +0.041$). This note does not touch that result; the sealed split remains evaluated exactly once.

1.2 What generality requires of a reader

Porting surfaced hard prerequisites that any candidate reader must satisfy for the method to even run:

1. **White-box access.** The method reads attention weights, edits the KV cache, and injects a 4-D mask. Any API / closed-weights model is categorically impossible.
2. **In-line visual tokens in the decoder’s self-attention KV.** Blocks are partitions of visual token *positions in the self-attention sequence*. A reader that injects vision via *cross-attention* (vision lives outside the self-attention KV) has no visual KV blocks to select or mask.
3. **Hookable eager attention returning 4-D weights, plus a 4-D additive mask honored by the language model** (the transferability crux; native TRANSFORMERS classes that delegate to a native decoder satisfy this, custom `trust_remote_code` attentions may not).
4. **Reconstructable, continuable positions** (1-D *arange* or M-RoPE) and enough context to prefill a multi-page document. Sub-native context ceilings force a reduced-budget setting (a disclosed confound, §4).
5. Empirically (§3), **the reader’s own attention must actually be a good below-page retriever**. Mechanical portability is necessary but not sufficient.

1.3 Protocol

All second-reader numbers are on the fixed **80-document / 444-query DEV pin** (`paq_confirm_pin.json`, MMLongBench-Doc DEV), reported as a *generality control*, not a new sealed claim. Each reader is compared to itself under two selectors at a matched budget: PAQ-KV

(own attention selects $\approx 5\%$ of blocks) vs. a *comparator* arm in which the same reader reads the top-5% pages chosen by cached ColQwen2 retrieval scores (no new ColQwen2 runs). We also record a *full-context* reference. The primary statistic is the doc-cluster 95% lower bound LB_{95} of (paqkv – comparator) (seed 0, 10,000 bootstrap resamples); the pre-registered “transfers” bar is $LB_{95} > 0$. An adversarial worst-case imputation of any missing query binds the verdict. A hard **identity gate** (teacher-forced argmax parity of the all-true 4-D mask vs. native causal decode, plus a restrictive mask changing the output, on ≥ 2 distinct documents including one ≥ 30 pages) must pass before any number is trusted. Verdicts are labeled INCOMPATIBLE (identity fails / port infeasible), INCOMPLETE (coverage/power short), `READER2_NO_TRANSFER` ($LB_{95} \leq 0$ with clean coverage), or `READER2_TRANSFERS` ($LB_{95} > 0$).

2 Results

2.1 Main comparison (verified readers)

Table 1 reports the three verified readers on the same 444-query DEV pin. Accuracy is the MMLongBench-Doc scorer; higher is better. All three readers read the same content selectors (own PAQ-KV blocks vs. ColQwen2 pages) at a matched $\approx 5\%$ budget.

- **Qwen3-VL-8B (native reader): win.** PAQ-KV – ColQwen2 = +0.076, LB_{95} = +0.041; PAQ-KV even edges its own full context (0.4042 vs. 0.3965), i.e. selection *denoises*.
- **Qwen3-VL-4B (same architecture, half the parameters): near-win / method holds.** Δ = +0.030, LB_{95} = –0.004 (just misses the bar at $n = 444$; the point estimate is a clear positive). PAQ-KV again *denoises* above its own full context (0.3640 vs. 0.3480), and on the 137 answerable queries it recovers **0.869** vs. the comparator’s **0.720**—the strongest below-page selection of any second reader. Per-query: 63 wins / 336 ties / 45 losses. Because the input geometry is byte-identical to the 8B (verified on all 80 documents), the gap to the 8B (–0.040 cross-reader) is a pure capacity effect, not a budget artifact. The method’s mechanism clearly survives shrinking the reader within the same architecture.
- **Qwen2.5-VL-7B (same Qwen-VL family): tie.** Δ = +0.0074, LB_{95} = [–0.037, +0.049]—a statistical dead heat straddling zero (PAQ-KV nominally ahead). Verdict `READER2_NO_TRANSFER` by the $LB_{95} > 0$ bar, but this is “no detectable difference,” not a loss. Per-query: 54 wins / 331 ties / 59 losses. On the 128 queries answerable by full context, PAQ-KV recovers **0.703** vs. the comparator’s **0.581**—here PAQ-KV is the *better* below-page selector. It loses only –0.054 [–0.087, –0.024] to its own full context.
- **InternVL3.5-8B (cross-family, Qwen3 LM): significant negative.** PAQ-KV – ColQwen2 = –0.067, LB_{95} = [–0.106, –0.029] (whole interval below zero). It also trails its own full context (–0.076 [–0.113, –0.041]). On answerable queries it recovers only **0.529**, losing $\approx 47\%$ of them to its 8% selection.

Verdict `READER2_NO_TRANSFER` (a genuine, well-powered negative).

2.2 Additional readers (in progress)

Four further readers were requested for this study. Qwen3-VL-4B is **complete** and appears in Tables 1 and 3 (it is the same TRANSFORMERS class as the sealed 8B, so it reused the audited harness with a loader swap and reproduced the 8B’s byte-identical input geometry). The remaining three (Table 2) each require a genuine per-model port because their attention APIs differ from the standard TRANSFORMERS 4-D-mask path the method relies on. Feasibility is decided cheaply by a Phase-0 identity spike before any full run; full ports proceed only for spike-passers, and a reader that fails the identity gate is reported honestly as INCOMPATIBLE rather than forced.

Per-reader notes.

- **Qwen3-VL-4B** is the same architecture as the sealed 8B (same TRANSFORMERS class, M-RoPE, native per-image budget), so the existing harness runs with a loader swap. It is a *within-architecture size ablation*, not a cross-architecture test; a win here would say the method survives shrinking the same reader, not that it transfers across families.
- **Molmo-7B-D** → INCOMPATIBLE (version drift, not mechanism). Its custom attention *does* expose the exact injection point the method needs—the top-level forward accepts a 4-D attention `_bias` fed straight to SDPA—so the mechanism is portable in principle. But its 2024 remote-code modeling file will not load under the pinned TRANSFORMERS 5.5.4: a cascade of tied-weights API mismatches (all `tied_weights_keys`, then `tie_weights(missing_keys=...)`) blocks from `_pretrained`. Loading it would require porting Molmo’s modeling file to the current API, out of scope here. This is an engineering/version-compatibility incompatibility, not a statement about the method.
- **MiniCPM-V-4.5** → INCOMPATIBLE (version drift, not mechanism). The model weights load, but its 2025 remote-code *processor* cannot build inputs under the pinned TRANSFORMERS 5.5.4: it expects a custom tokenizer attribute (`tokenizer.im_start_id`) that the current fast-tokenizer backend does not expose. Independently, its 3D-Resampler compresses each page to ≈ 64 tokens, so even a working port would be near-degenerate below page granularity (a 5% block budget of ≈ 64 tokens is < 1 block per page); and its LM is Qwen3, the same family as the InternVL negative.
- **PaliGemma-2** → INCOMPATIBLE (representation, not version). It loads natively and accepts multiple images (1024 tokens/page), but its Gemma-2 language model has an ≈ 8192 -token context with a 4096 sliding window, and PaliGemma uses a *fixed* 1024 tokens per image with no dynamic per-page budget. That caps a document at ≈ 7 pages; the study’s documents are 17–104 pages (median ≈ 32), so PAQ-KV’s “prefill the whole document once” premise cannot be satisfied—the document does not fit.

Table 1: PAQ-KV vs. the ColQwen2-page comparator, within each reader, on the 80-doc / 444-q DEV pin. Δ = mean(paqkv – ColQwen2); LB_{95} = doc-cluster 95% lower bound (pre-registered “transfers” bar is $LB_{95} > 0$). “ctx” is the mean realized prefill length. Qwen3-VL-8B is the reader the method was developed and sealed on; Qwen3-VL-4B uses byte-identical input geometry (a pure size ablation).

Reader	ctx	PAQ-KV	ColQwen2	Full	Δ	LB_{95}
Qwen3-VL-8B (<i>native, sealed</i>)	~69K	0.4042	0.3284	0.3965	+0.076	+0.041
Qwen3-VL-4B (<i>size ablation</i>)	~69K	0.3640	0.3336	0.3480	+0.030	−0.004
Qwen2.5-VL-7B	~52K	0.2698	0.2623	0.3237	+0.007	−0.037
InternVL3.5-8B	~36K	0.1876	0.2542	0.2637	−0.067	−0.106

Table 2: Additional candidate readers. “Native” = a TRANSFORMERS 5.5 class exists (mask machinery and eager hooks transfer as for InternVL/Qwen2.5-VL); “remote” = trust_remote_code custom attention (4-D mask honored and hookability must be validated by the spike). Results are filled in as runs complete.

Reader	LM backbone	Vision integration	Impl.	Phase-0 verdict
Molmo-7B-D	Qwen2-7B	CLIP, in-line tokens	remote	INCOMPATIBLE (model won’t load: version drift)
MiniCPM-V-4.5	Qwen3-8B	3D-Resampler (64 tok)	remote	INCOMPATIBLE (processor won’t build inputs)
PaliGemma-2	Gemma-2	prefix-LM (bidir.)	native	INCOMPATIBLE (~8K ctx: cannot fit a doc)

(Its prefix-LM attention and short-output tuning would further complicate a port even at fitting lengths.)

Takeaway from the three incompatibilities. They are not method failures; they map the method’s *practical* deployment envelope. Running reader-own PAQ-KV needs, in addition to the mechanistic property of §3: (i) a version-current implementation exposing a hookable attention and an honored additive mask; and (ii) enough context, with a dynamic per-page token budget, to prefill a whole multi-page document. Among the readers surveyed, the Qwen-VL family (Qwen3-VL-8B/4B, Qwen2.5-VL-7B) and InternVL3.5-8B satisfy both; the three above each miss one. This is worth stating plainly as a scope condition rather than omitting it.

3 Discussion: a mechanistic generality gradient

The verified readers order cleanly by architectural distance from the reader the method was developed on (Table 3). The benefit is not universal; it *degrades monotonically* as the reader moves away from Qwen3-VL-8B, and the degradation tracks a measurable property—how well the reader’s own question-to-context attention localizes the answer below page granularity. Qwen3-VL-8B has this property (it recovers answerable queries well at 5% and even denoises above full context); **Qwen3-VL-4B retains it** (near-win at matched, byte-identical budget; denoises above full; recovers 0.87 of answerable queries), so the effect survives halving the parameter count *within the same architecture* and is not an artifact of the 8B’s scale alone; Qwen2.5-VL-7B partially has it (break-even; recovers 0.70); InternVL3.5-8B lacks it here (recovers 0.53 and loses). The two Qwen3-VL points (win, near-win) versus the same-family tie and the cross-family negative locate the boundary at *architecture*, not merely scale. This *confirms* rather than refutes the single-reader observation, and reframes it from an unexamined weakness into an investigated boundary condition:

PAQ-KV is a white-box method whose benefit is contingent on the host reader’s attention being an effective below-page retriever.

3.1 Can parameter tuning close the InternVL gap? (DEV-only, pre-registered, held-out)

A natural follow-up: is InternVL’s deficit an artifact of the fixed 5% budget or the default attention read-out, and can tuning close it? We answered this under a frozen pre-registration that never touches the sealed set: a TUNE(40)/HELDOUT(40) document split of the DEV pin (seed 0, stratified, positive-allowlist enforced), a frozen grid over budget $R \in \{5, 10, 20, 30\}\%$, attention read-out {mean-all, last-quarter-layers, max-over-heads, entropy-weighted}, token→block aggregation, and a boundary-keep policy (measured in *attended visual tokens*), with the winning configuration selected on TUNE and evaluated once on HELDOUT. A pre-GPU power analysis fixed the reading: at $n=40$ documents the doc-cluster half-width is ~ 0.06 , so “matching” cannot register as superiority—hence the primary is the descriptive budget frontier (all 80 docs) plus the low-variance paired C^* —default contrast and a noninferiority margin $\delta=0.03$.

Budget is the only lever, and it does not yield a win.

The frontier (Table 4) shows InternVL’s gap to its comparator narrowing monotonically with budget (−0.067 at 5% to −0.001 at 30%), and the method reaching its *own full context* by $\sim 20\%$ (recovery of answerable queries rises from 0.53 at 5% to 0.74 at 10%). But LB_{95} (method – comparator) stays below zero at *every* budget—it never becomes a significant win. The tuned winner C^* (max-over-heads read-out, $R^*=10\%$) gives **no gain beyond budget** on HELDOUT (paired C^* —default: $\Delta+0.016$, $LB_{95}-0.014$), and is **not noninferior** to the comparator at R^* ($\Delta-0.006$, $LB_{95}-0.055 < -\delta$), despite the method using *more* attended visual tokens than the comparator there (2126 vs. 1768). So a smarter attention read-out does not rescue In-

Table 3: The generality gradient, ordered by architectural distance from the sealed reader. Δ / doc-cluster LB_{95} (paqkv – comparator), DEV pin (Qwen3-VL-8B also shows a sealed out-of-sample $LB_{95} = +0.029$). The benefit degrades monotonically with distance.

Reader	Relation to the sealed reader	Δ / LB_{95}
Qwen3-VL-8B	native (sealed reader)	+0.076 / +0.041 (<i>win</i>)
Qwen3-VL-4B	same architecture, 4B	+0.030 / −0.004 (<i>near-win</i>)
Qwen2.5-VL-7B	same family, prior gen.	+0.007 / −0.037 (<i>tie</i>)
InternVL3.5-8B	cross-family VLM (Qwen3 LM)	−0.067 / −0.106 (<i>neg.</i>)

ternVL: only spending several times the budget narrows the within-reader gap, and even then only to break-even with its own full context, never to a comparator win. This *reinforces* §3: InternVL’s own attention is a genuinely weaker below-page retriever, and this cannot be tuned away with these knobs. (All values are DEV held-out estimates, pre-registered; the sealed set is untouched and no sealed claim is made.)

4 Confounds and honesty caveats

Any use of these numbers must carry the following.

1. **Reduced-budget confound.** Both second readers run below the native-long 8B budget (InternVL $\sim 36K$ via a 40K context ceiling; Qwen2.5-VL $\sim 52K$ via a 128K ceiling plus a processor pre-resize). Within-reader matched-budget comparisons carry the findings; cross-reader comparisons *against* the 8B are confounded by capacity and budget and are secondary only.
2. **Shared LM family.** InternVL3.5-8B’s language backbone is Qwen3, which is *why* the KV/attention/mask machinery transfers. It is a legitimate cross-VLM result but not a fully independent language architecture; disclose it.
3. **Conservative-against-method budget asymmetry (Qwen2.5-VL).** The method downscales all pages in the full-document prefill, whereas the comparator renders its 1–5 kept pages near-native and thus gets a slightly larger per-page budget. The tie is therefore a floor on the within-family effect.
4. **DEV-pin, not a new sealed claim.** All second-reader numbers are on the DEV pin as a generality control. The sealed split is untouched.

Verification (Appendix A)

Every reported run passed a hard identity gate and was independently checked. For each verified reader: the identity audit confirmed the all-true 4-D mask reproduces native causal decoding under the teacher-forced argmax criterion ($\max |\Delta \text{logit}| = 0$), a restrictive mask changes the output, on ≥ 2 distinct documents (one ≥ 30 pages), with the resolved model revision and code hash bound into the frozen run manifest. A CPU sparse-mask audit confirmed the block-to-column mapping is exact (all-blocks selection reproduces the all-true mask; a sparse subset masks exactly the intended columns). An off-meter code reviewer independently attempted to falsify each result and found no implementation defect (“genuine negative”

for InternVL; “genuine null with favorable caveats” for Qwen2.5-VL). Full coverage for both second readers: 444/444 queries, 80 documents, 0 skipped, 0 dropped, identity pass. Artifacts, verdict JSONs, reviewer transcripts, and audit scripts are in the project repository under `research/` and `research/reviews/paq_second_vlm_2026-07-22/`; the consolidated running log is `research/design_notes/paq_cross_reader_generality_2026-07-23.md`.

References (model / benchmark sources)

- Qwen3-VL / Qwen2.5-VL** Qwen Team. <https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct>
- InternVL3.5-8B** OpenGVLab. https://huggingface.co/OpenGVLab/InternVL3_5-8B-HF
- Molmo-7B-D** Allen Institute for AI. <https://huggingface.co/allenai/Molmo-7B-D-0924>
- MiniCPM-V-4.5** OpenBMB. https://huggingface.co/openbmb/MiniCPM-V-4_5
- PaliGemma-2** Google. <https://huggingface.co/google/paligemma2-3b-mix-448>
- ColQwen2** Faysse et al. <https://huggingface.co/vidore/colqwen2-v1.0>
- MMLongBench-Doc** Ma et al. <https://mmlongbench-doc.github.io>

Table 4: InternVL3.5-8B budget frontier on all 80 DEV documents (default read-out; descriptive, no selection). Δ , LB_{95} are doc-cluster (method — comparator). The gap narrows with budget but LB_{95} never clears zero; the method reaches its own full context by $\sim 20\%$.

R	method	comparator	full	$\Delta(\text{m}-\text{c})$	$\text{LB}_{95}(\text{m}-\text{c})$
5%	0.1876	0.2542	0.2637	−0.067	−0.106
10%	0.2367	0.2762	0.2637	−0.040	−0.083
20%	0.2664	0.2767	0.2637	−0.010	−0.048
30%	0.2684	0.2691	0.2637	−0.001	−0.036